

Managing many models

February 2017

Hadley Wickham
[@hadleywickham](#)
Chief Scientist, RStudio

You've never seen data presented like this. With the drama and urgency of a sportscaster, statistics guru Hans Rosling debunks myths about the so-called "developing world."

https://www.ted.com/talks/hans_rosling_shows_the_best_stats_you_ve_ever_seen

Chart

Map



How to use



Share graph



Normal view

Color

Regions of the world

Geographic regions



Select

- ☐ Afghanistan
- ☐ Albania
- ☐ Algeria
- ☐ Andorra
- ☐ Angola
- ☐ Antigua and Barb...
- ☐ Argentina
- ☐ Armenia
- ☐ Aruba
- ☐ Australia
- ☐ Austria
- ☐ Azerbaijan
- ☐ Bahamas

☐ Deselect all

Size

Sources(s)

Population, total



Life expectancy (years)

lin

lin

Source(s)

85
80
75
70
65
60
55
50
45
40
35
30
25

200 400 1 000 2 000 4 000 10 000 20 000 40 000

Income per person (GDP/capita, PPP\$ inflation-adjusted)

log

Stop

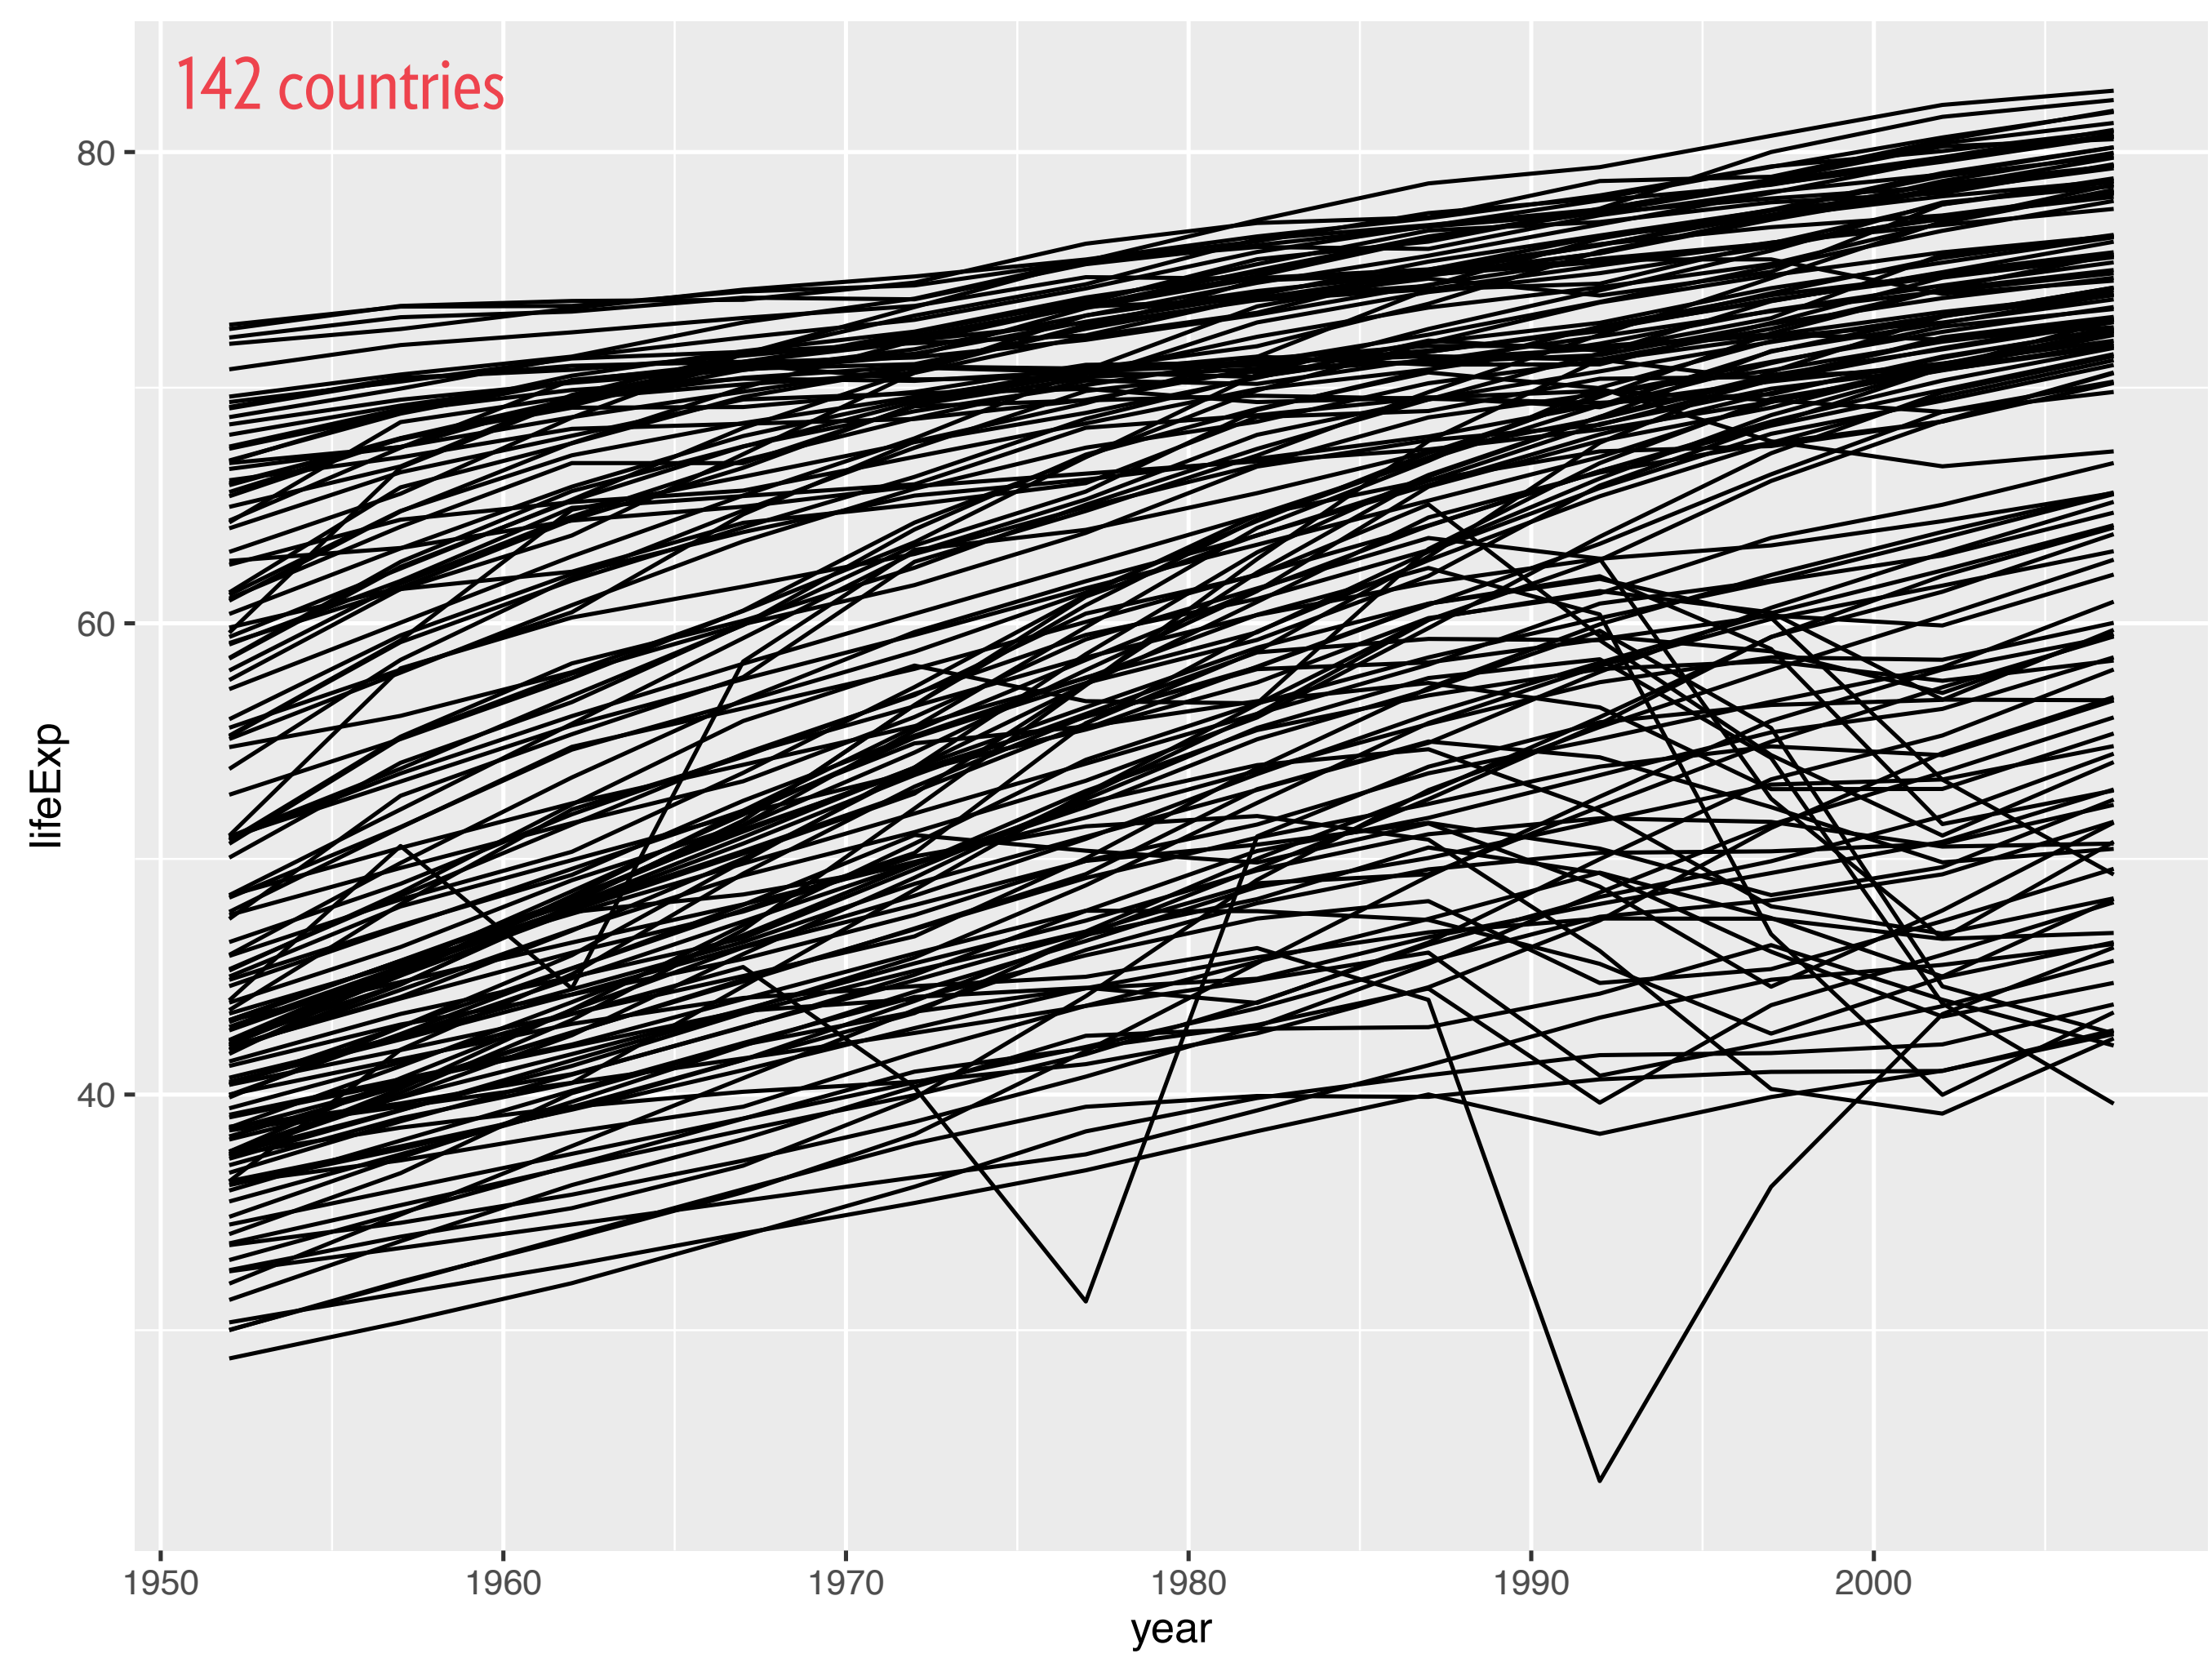
1800 1820 1840 1860 1880 1900 1920 1940 1960 1980 2000


☒ Trails

[Terms of use](#)

© Google 2008

142 countries





But...

Arbitrarily complicated models

Scales to **big** data

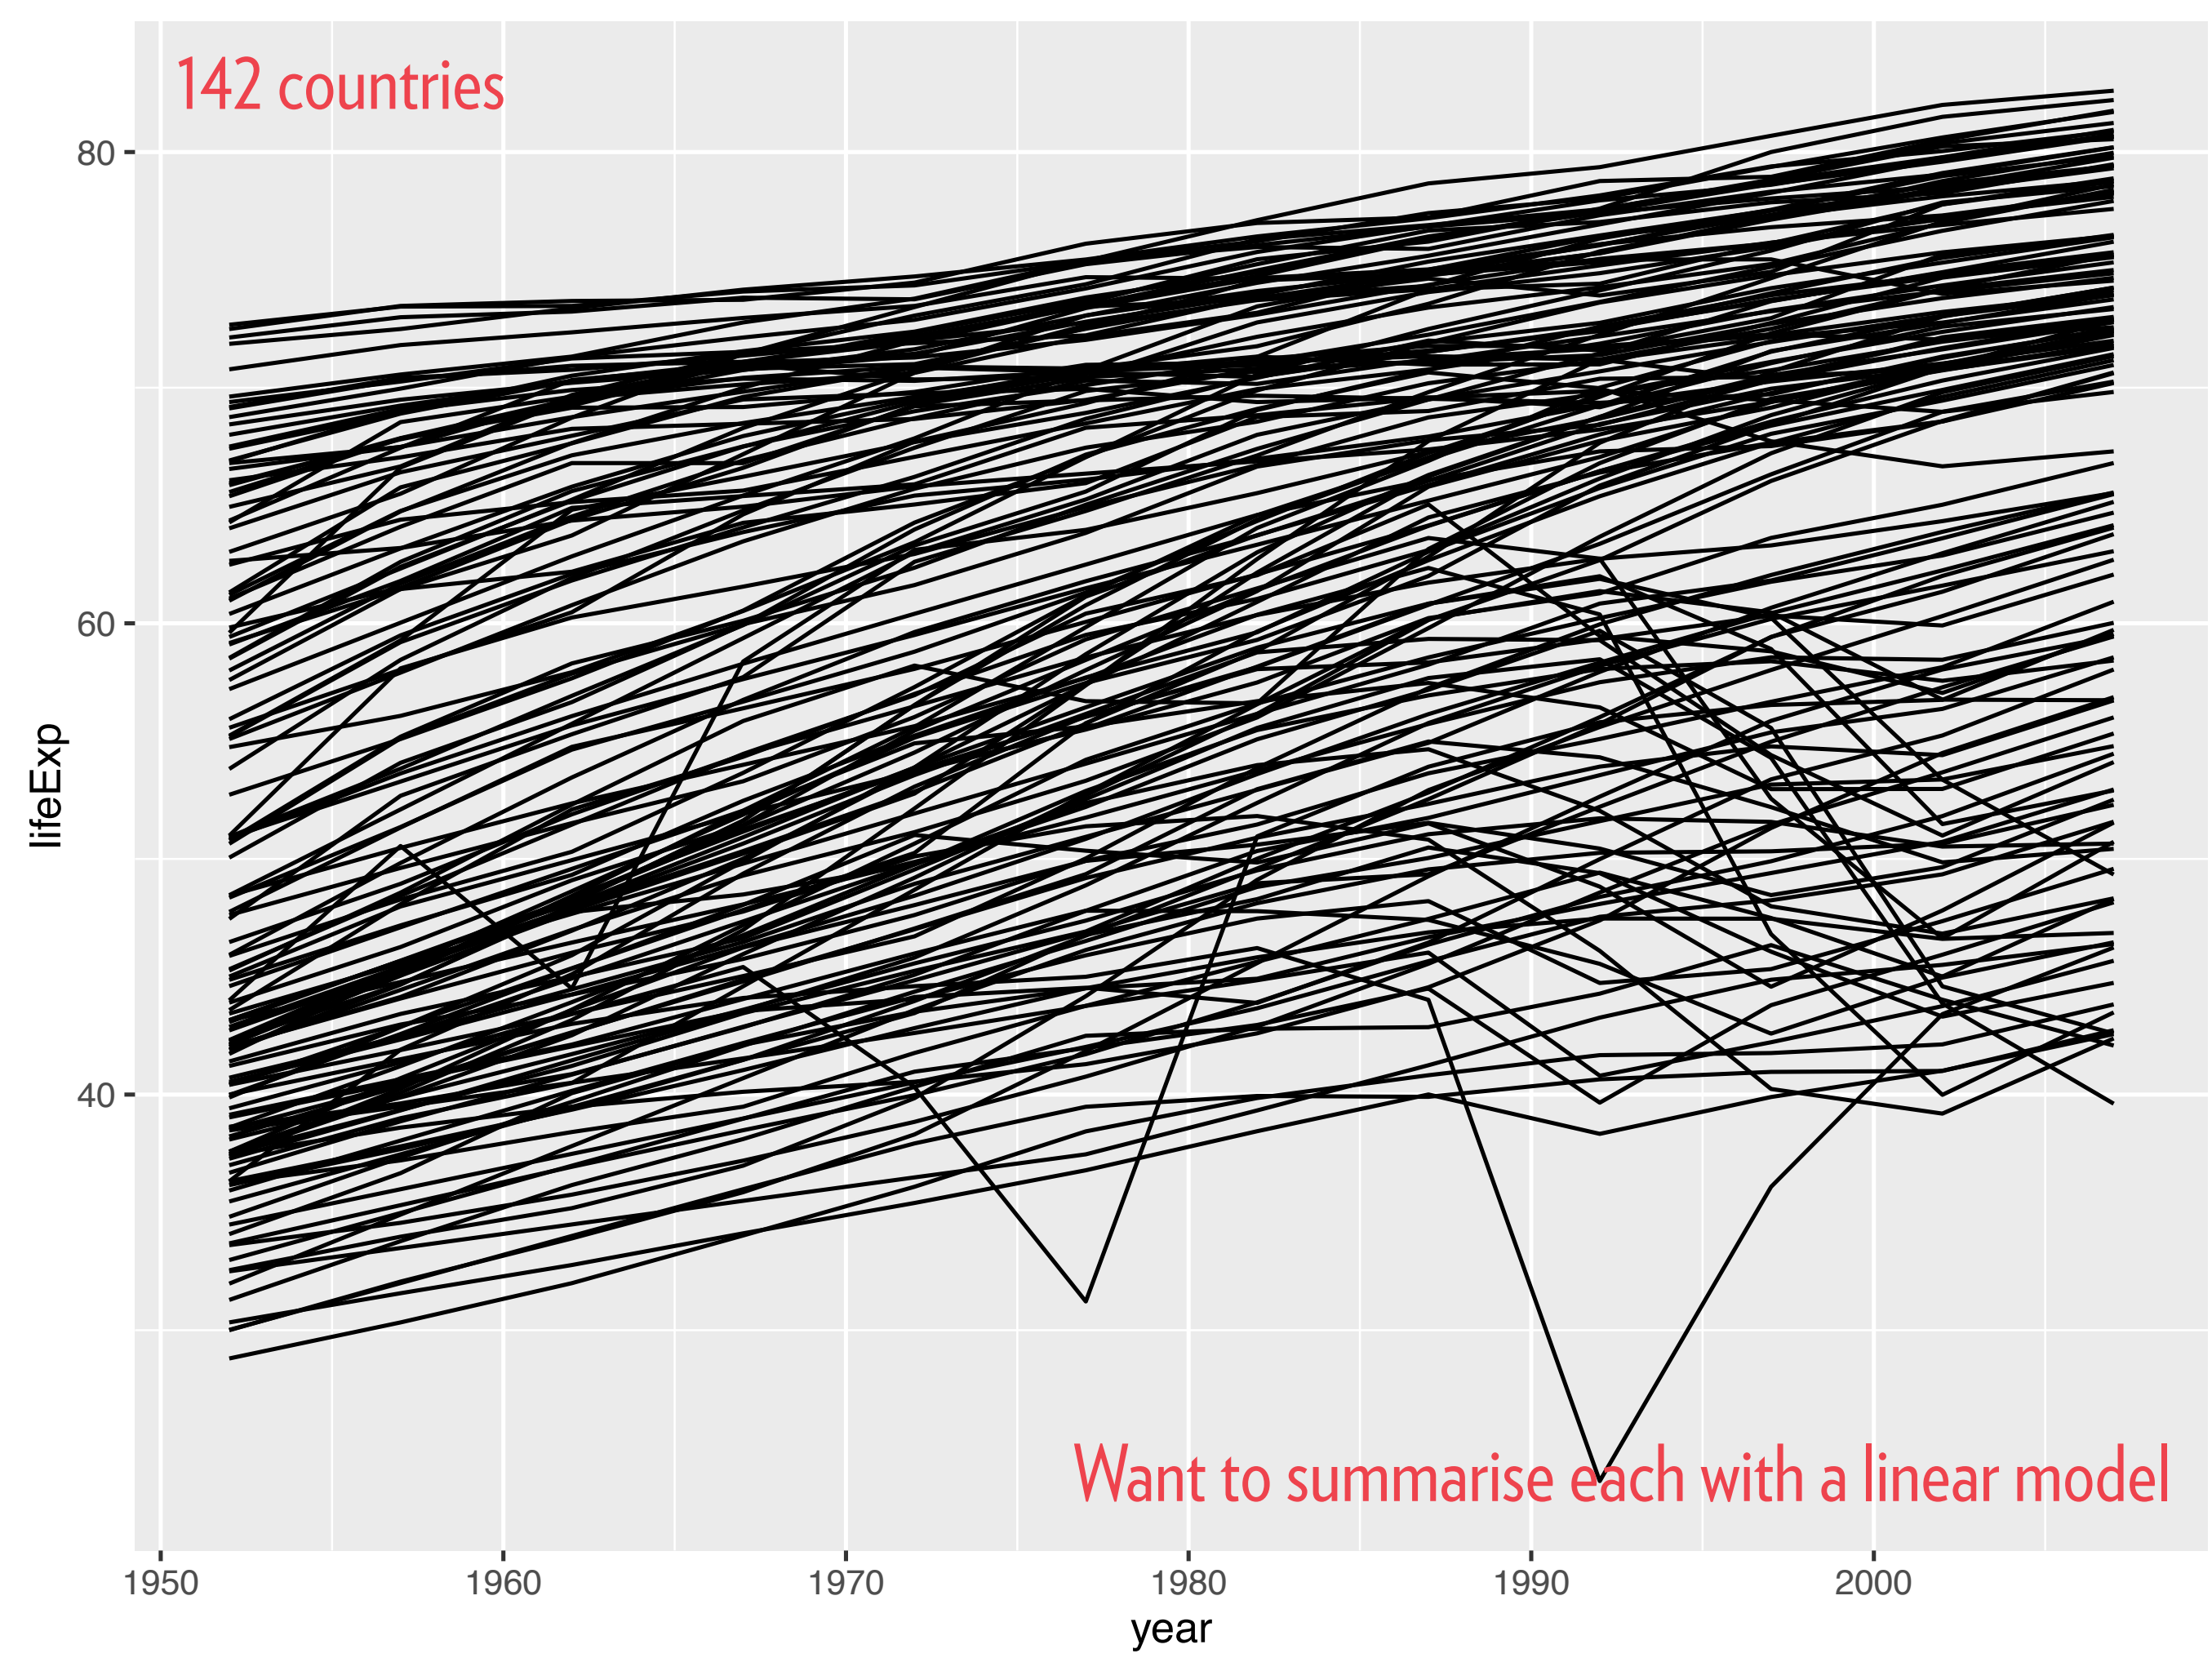
Three simple underlying ideas

Each idea is partnered with a package

1. Nested data (tidyr)
2. Functional programming (purrr)
3. Models \rightarrow tidy data (broom)

Nested data

142 countries

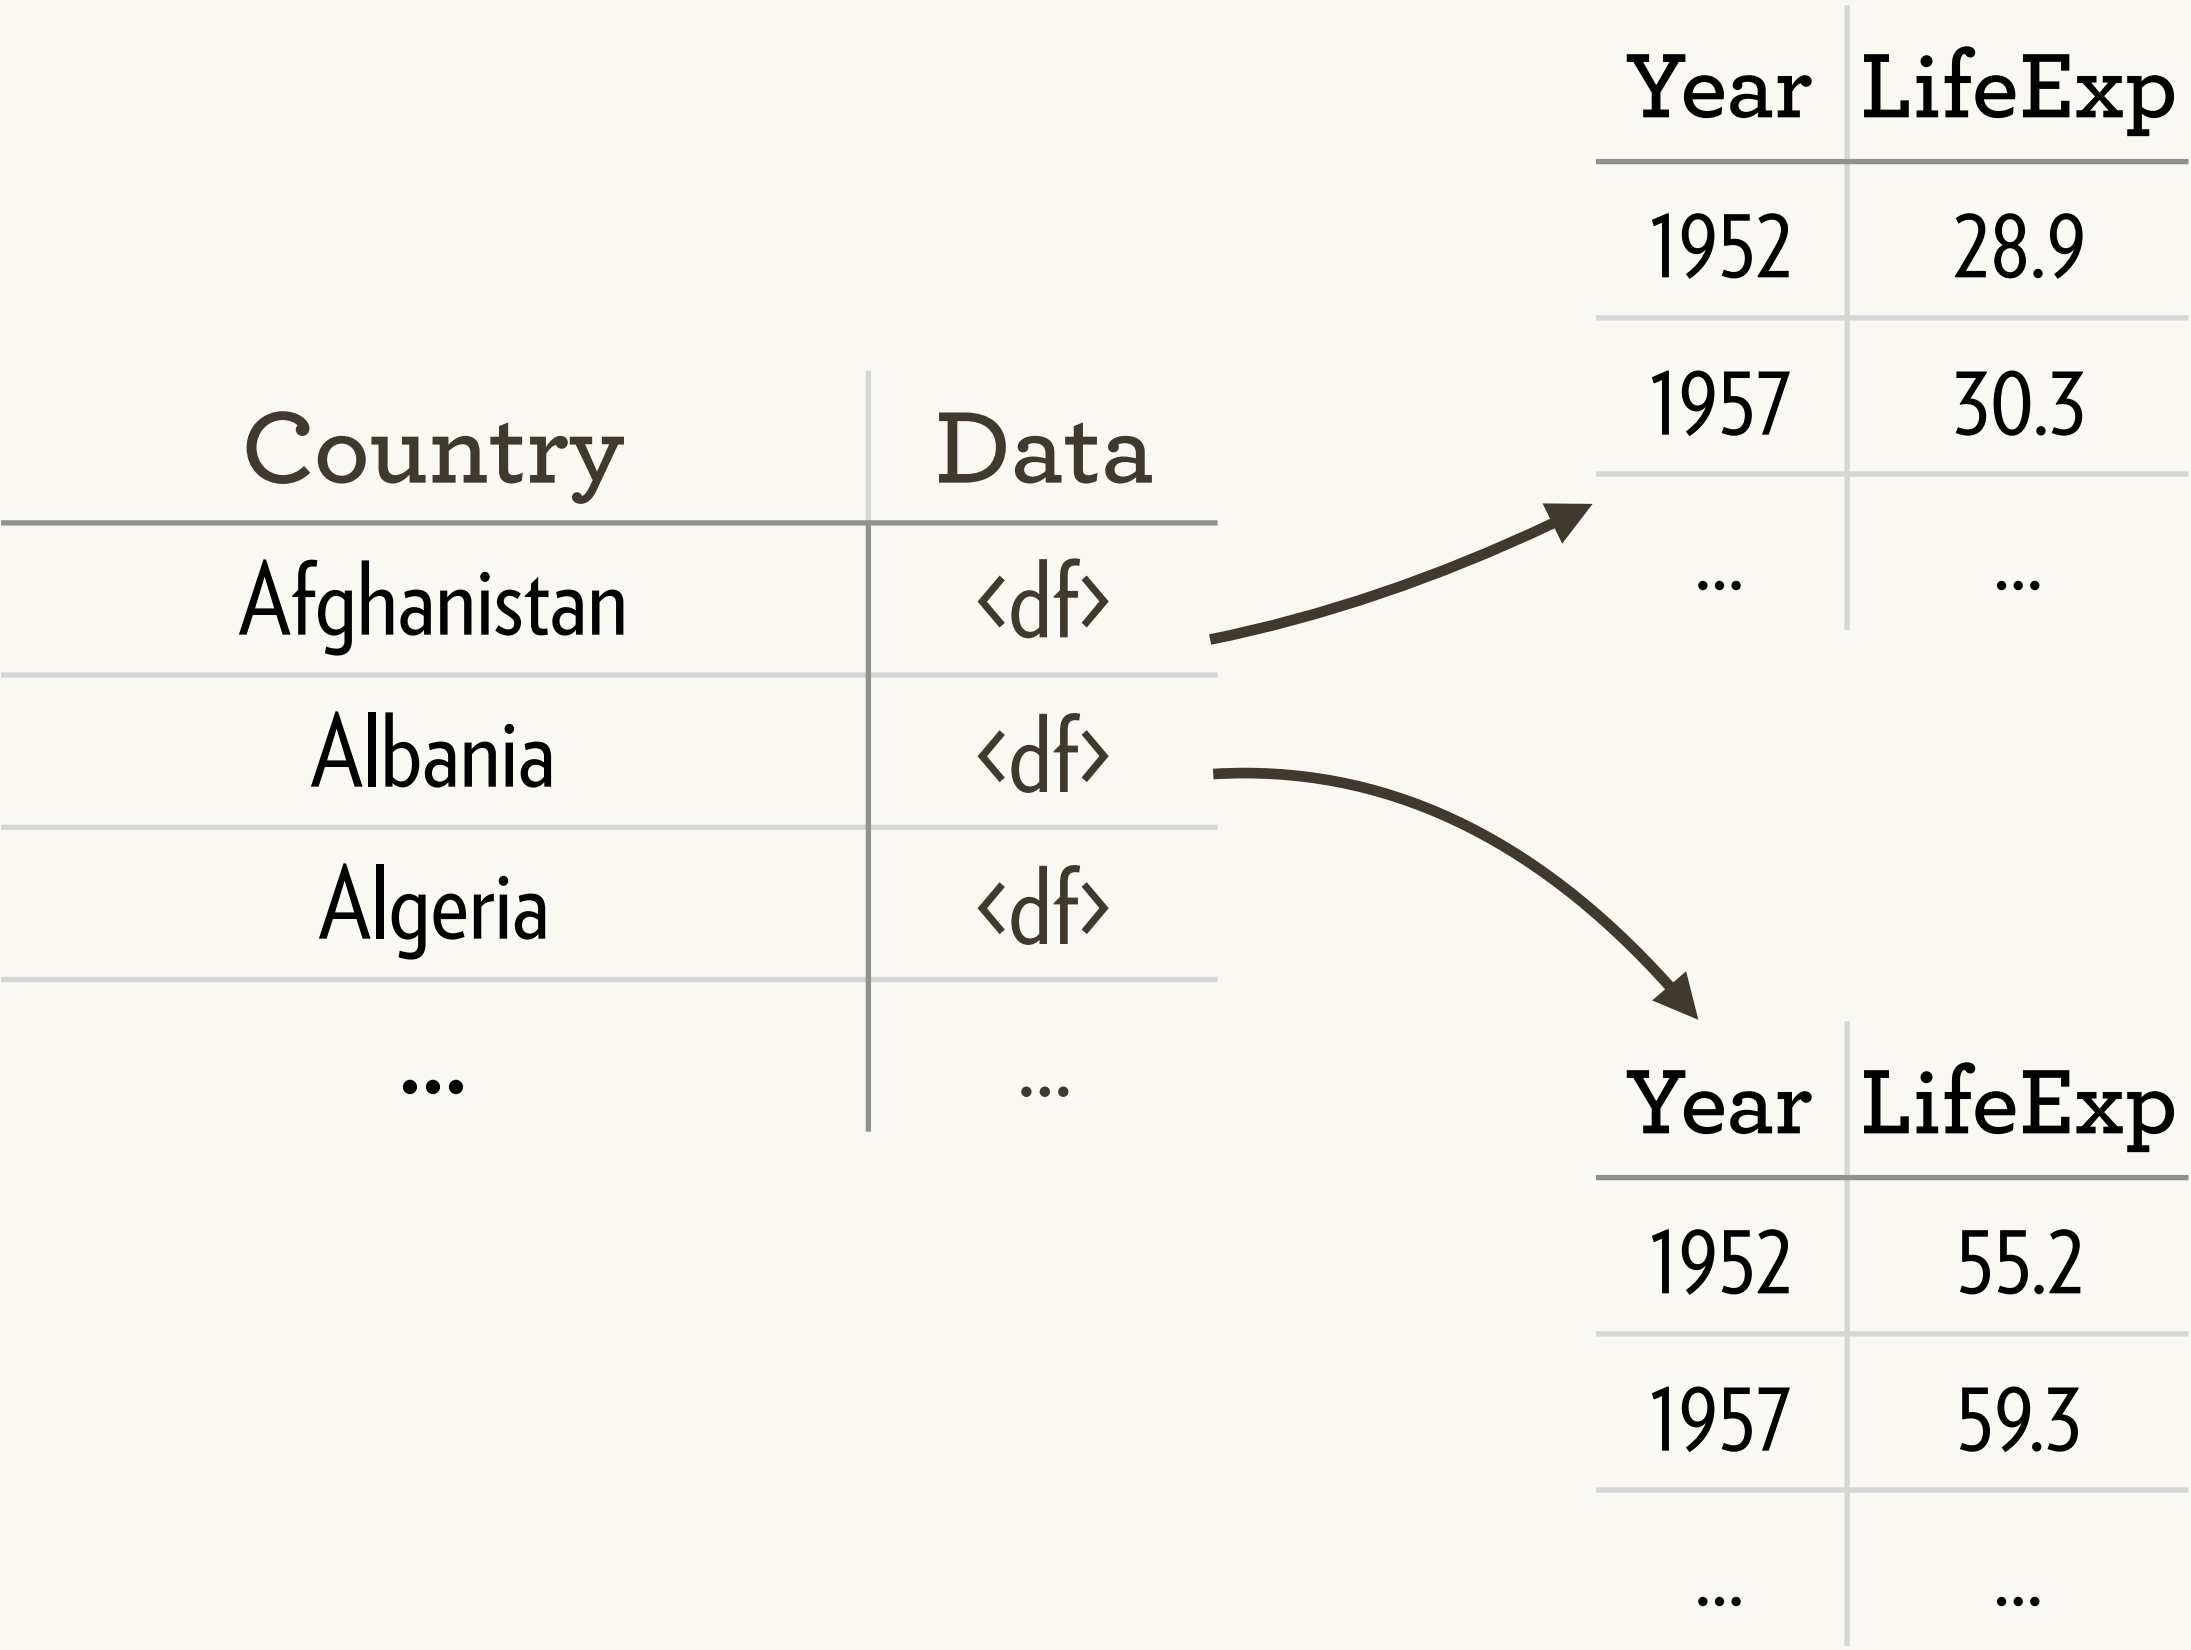


Want to summarise each with a linear model

Currently our data has one row per observation

Country	Year	LifeEx
Afghanistan	1952	28.9
Afghanistan	1957	30.3
Afghanistan
Albania	1952	55.2
Albania	1957	59.3
Albania
Algeria
...

More convenient to one row per group



I call this a **nested data frame**

In R:

```
library(dplyr)
```

```
library(tidyr)
```

```
by_country <- gapminder %>%
```

```
  group_by(continent, country) %>%
```

```
  nest()
```

Haven't seen pipes?

```
x %>% f(y)
```

```
# is the same as:
```

```
f(x, y)
```

```
gapminder %>%
```

```
  group_by(continent, country) %>%
```

```
  nest()
```

```
# same as:
```

```
nest(group_by(gapminder, continent, country))
```

Each country will have an associated model

Country	Data
Afghanistan	<df>
Albania	<df>
Algeria	<df>
...	...


`lm(lifeExp ~ year1950, data = afghanistan)`

`lm(lifeExp1950 ~ year, data = albania)`

Why not store that in a column too?

Country	Data	Model
Afghanistan	<df>	<lm>
Albania	<df>	<lm>
Algeria	<df>	<lm>
...

List-columns keep related things together



Anything can go in a list & a list can go in a data frame

In R:

```
library(dplyr)
```

```
library(purrr)
```

```
country_model <- function(df) {  
  lm(lifeExp ~ year1950, data = df)  
}
```

```
models <- by_country %>%  
  mutate(  
    mod = map(data, country_model)  
  )
```

Functional programming

Or, why for loops are “bad”

Motivated by baking cupcakes

Vanilla cupcakes

The hummingbird
bakery cookbook

1 cup flour
a scant $\frac{3}{4}$ cup sugar
1 $\frac{1}{2}$ t baking powder
3 T unsalted butter
 $\frac{1}{2}$ cup whole milk
1 egg
 $\frac{1}{4}$ t pure vanilla extract

Preheat oven to 350°F.

Put the flour, sugar, baking powder, salt, and butter in a freestanding electric mixer with a paddle attachment and beat on slow speed until you get a sandy consistency and everything is combined.

Whisk the milk, egg, and vanilla together in a pitcher, then slowly pour about half into the flour mixture, beat to combine, and turn the mixer up to high speed to get rid of any lumps.

Turn the mixer down to a slower speed and slowly pour in the remaining milk mixture. Continue mixing for a couple of more minutes until the batter is smooth but do not overmix.

Spoon the batter into paper cases until $\frac{2}{3}$ full and bake in the preheated oven for 20-25 minutes, or until the cake bounces back when touched.

Chocolate cupcakes

The hummingbird
bakery cookbook

$\frac{3}{4}$ cup + 2T flour
2 $\frac{1}{2}$ T cocoa powder
a scant $\frac{3}{4}$ cup sugar
1 $\frac{1}{2}$ t baking powder
3 T unsalted butter
 $\frac{1}{2}$ cup whole milk
1 egg
 $\frac{1}{4}$ t pure vanilla extract

Preheat oven to 350°F.

Put the flour, cocoa, sugar, baking powder, salt, and butter in a freestanding electric mixer with a paddle attachment and beat on slow speed until you get a sandy consistency and everything is combined.

Whisk the milk, egg, and vanilla together in a pitcher, then slowly pour about half into the flour mixture, beat to combine, and turn the mixer up to high speed to get rid of any lumps.

Turn the mixer down to a slower speed and slowly pour in the remaining milk mixture. Continue mixing for a couple of more minutes until the batter is smooth but do not overmix.

Spoon the batter into paper cases until $\frac{2}{3}$ full and bake in the preheated oven for 20-25 minutes, or until the cake bounces back when touched.

Chocolate cupcakes

The hummingbird
bakery cookbook

$\frac{3}{4}$ cup + 2T flour

2 $\frac{1}{2}$ T cocoa powder

a scant $\frac{3}{4}$ cup sugar

1 $\frac{1}{2}$ t baking powder

3 T unsalted butter

$\frac{1}{2}$ cup whole milk

1 egg

$\frac{1}{4}$ t pure vanilla extract

Preheat oven to 350°F.

Put the flour, **cocoa**, sugar, baking powder, salt, and butter in a freestanding electric mixer with a paddle attachment and beat on slow speed until you get a sandy consistency and everything is combined.

Whisk the milk, egg, and vanilla together in a pitcher, then slowly pour about half into the flour mixture, beat to combine, and turn the mixer up to high speed to get rid of any lumps.

Turn the mixer down to a slower speed and slowly pour in the remaining milk mixture. Continue mixing for a couple of more minutes until the batter is smooth but do not overmix.

Spoon the batter into paper cases until $\frac{2}{3}$ full and bake in the preheated oven for 20-25 minutes, or until the cake bounces back when touched.

Vanilla cupcakes

The hummingbird
bakery cookbook

1 cup flour
a scant $\frac{3}{4}$ cup sugar
1 $\frac{1}{2}$ t baking powder
3 T unsalted butter
 $\frac{1}{2}$ cup whole milk
1 egg
 $\frac{1}{4}$ t pure vanilla extract

Preheat oven to 350°F.

Put the flour, sugar, baking powder, salt, and butter in a freestanding electric mixer with a paddle attachment and beat on slow speed until you get a sandy consistency and everything is combined.

Whisk the milk, egg, and vanilla together in a pitcher, then slowly pour about half into the flour mixture, beat to combine, and turn the mixer up to high speed to get rid of any lumps.

Turn the mixer down to a slower speed and slowly pour in the remaining milk mixture. Continue mixing for a couple of more minutes until the batter is smooth but do not overmix.

Spoon the batter into paper cases until $\frac{2}{3}$ full and bake in the preheated oven for 20-25 minutes, or until the cake bounces back when touched.

Vanilla cupcakes

The hummingbird
bakery cookbook

120g flour

140g sugar

1.5 t baking powder

40g unsalted butter

120ml milk

1 egg

0.25 t pure vanilla extract

Preheat oven to 170°C.

Put the flour, sugar, baking powder, salt, and butter in a freestanding electric mixer with a paddle attachment and beat on slow speed until you get a sandy consistency and everything is combined.

Whisk the milk, egg, and vanilla together in a pitcher, then slowly pour about half into the flour mixture, beat to combine, and turn the mixer up to high speed to get rid of any lumps.

Turn the mixer down to a slower speed and slowly pour in the remaining milk mixture. Continue mixing for a couple of more minutes until the batter is smooth but do not overmix.

Spoon the batter into paper cases until 2/3 full and bake in the preheated oven for 20-25 minutes, or until the cake bounces back when touched.

Vanilla cupcakes

The hummingbird
bakery cookbook

120g flour
140g sugar
1.5 t baking powder
40g butter
120ml milk
1 egg
0.25 t vanilla

Beat flour, sugar, baking powder, salt, and butter until sandy.

Whisk milk, egg, and vanilla. Mix half into flour mixture until smooth (use high speed). Beat in remaining half. Mix until smooth.

Bake 20-25 min at 170°C.

Vanilla cupcakes

The hummingbird
bakery cookbook

120g flour

140g sugar

1.5 t baking powder

40g butter

120ml milk

1 egg

0.25 t vanilla

Beat **dry ingredients** + butter until sandy.

Whisk together **wet ingredients**. Mix half into dry until smooth (use high speed). Beat in remaining half. Mix until smooth.

Bake 20-25 min at 170°C.

Cupcakes

Beat dry ingredients + butter until sandy.

Whisk together wet ingredients. Mix half into dry until smooth (use high speed). Beat in remaining half. Mix until smooth.

Bake 20-25 min at 170°C.

Vanilla

120g flour

140g sugar

1.5t baking powder

40g butter

120ml milk

1 egg

0.25 t vanilla

Chocolate

100g flour

20g cocoa

140g sugar

1.5t baking powder

40g butter

120ml milk

1 egg

0.25 t vanilla

Cupcakes

	Flour	Baking powder	Sugar	Butter	Egg	Extra
Vanilla	120	1.5	140	40	1	0.25t vanilla
Chocolate	100	1.5	140	40	1	20g cocoa • 0.25t vanilla
Lemon	120	1.5	140	40	1	2T lemon zest
Red velvet	150	0	150	60	1	10g cocoa • 20ml red colouring • 1.5t vinegar • 0.5 t baking soda

For loops emphasise the objects

```
out1 <- vector("double", ncol(mtcars))
for(i in seq_along(mtcars)) {
  out1[[i]] <- mean(mtcars[[i]], na.rm = TRUE)
}
```

```
out2 <- vector("double", ncol(mtcars))
for(i in seq_along(mtcars)) {
  out2[[i]] <- median(mtcars[[i]], na.rm = TRUE)
}
```

For loops emphasise the objects

```
out1 <- vector("double", ncol(mtcars))  
for(i in seq_along(mtcars)) {  
  out1[[i]] <- mean(mtcars[[i]], na.rm = TRUE)  
}
```

```
out2 <- vector("double", ncol(mtcars))  
for(i in seq_along(mtcars)) {  
  out2[[i]] <- median(mtcars[[i]], na.rm = TRUE)  
}
```

Not the actions

```
out1 <- vector("double", ncol(mtcars))  
for(i in seq_along(mtcars)) {  
  out1[[i]] <- mean(mtcars[[i]], na.rm = TRUE)  
}
```

```
out2 <- vector("double", ncol(mtcars))  
for(i in seq_along(mtcars)) {  
  out2[[i]] <- median(mtcars[[i]], na.rm = TRUE)  
}
```

Functional programming emphasises the actions

```
library(purrr)
means <- mtcars %>% map_dbl(mean)
medians <- mtcars %>% map_dbl(median)
```

What does `map_dbl()` look like?

```
map_dbl <- function(x, f, ...) {  
  out <- vector("double", length(x))  
  for (i in seq_along(out)) {  
    out[i] <- f(x[[i]], ...)  
  }  
  out  
}
```

Actual implementation a little different

There are many variants:

```
map_int <- function(x, f, ...) {  
  out <- vector("integer", length(x))  
  for (i in seq_along(out)) {  
    out[i] <- f(x[[i]], ...)  
  }  
  out  
}
```

Some vary the output

```
map <- function(x, f, ...) {  
  out <- vector("list", length(x))  
  for (i in seq_along(out)) {  
    out[[i]] <- f(x[[i]], ...)  
  }  
  out  
}
```

This is the same as `lapply()`!

Others vary the input

```
map2 <- function(x, y, f, ...) {  
  out <- vector("list", length(x))  
  for (i in seq_along(out)) {  
    out[[i]] <- f(x[[i]], y[[i]], ...)  
  }  
  out  
}
```

We can even think of functions as data


```
funcs <- list(mean, median, sd)
```


```
funcs %>%
```

```
  map(~ mtcars %>% map_dbl(.x))
```

Back to gapminder

142 countries





We nested the data to get a list of data frames

Country	Year	LifeEx
Afghanistan	1952	28.9
Afghanistan	1957	30.3
Afghanistan
Albania	1952	55.2
Albania	1957	59.3
Albania
Algeria
...

`nest()`



Country	Data
Afghanistan	<df>
Albania	<df>
Algeria	<df>
...	...

Then we fitted a model to each country

```
library(dplyr)
library(tidyr)
library(purrr)

country_model <- function(df) {
  lm(lifeExp ~ year1950, data = df)
}

gapminder %>%
  group_by(continent, country) %>%
  nest() %>%
  mutate(
    mod = data %>% map(country_model)
  )
```

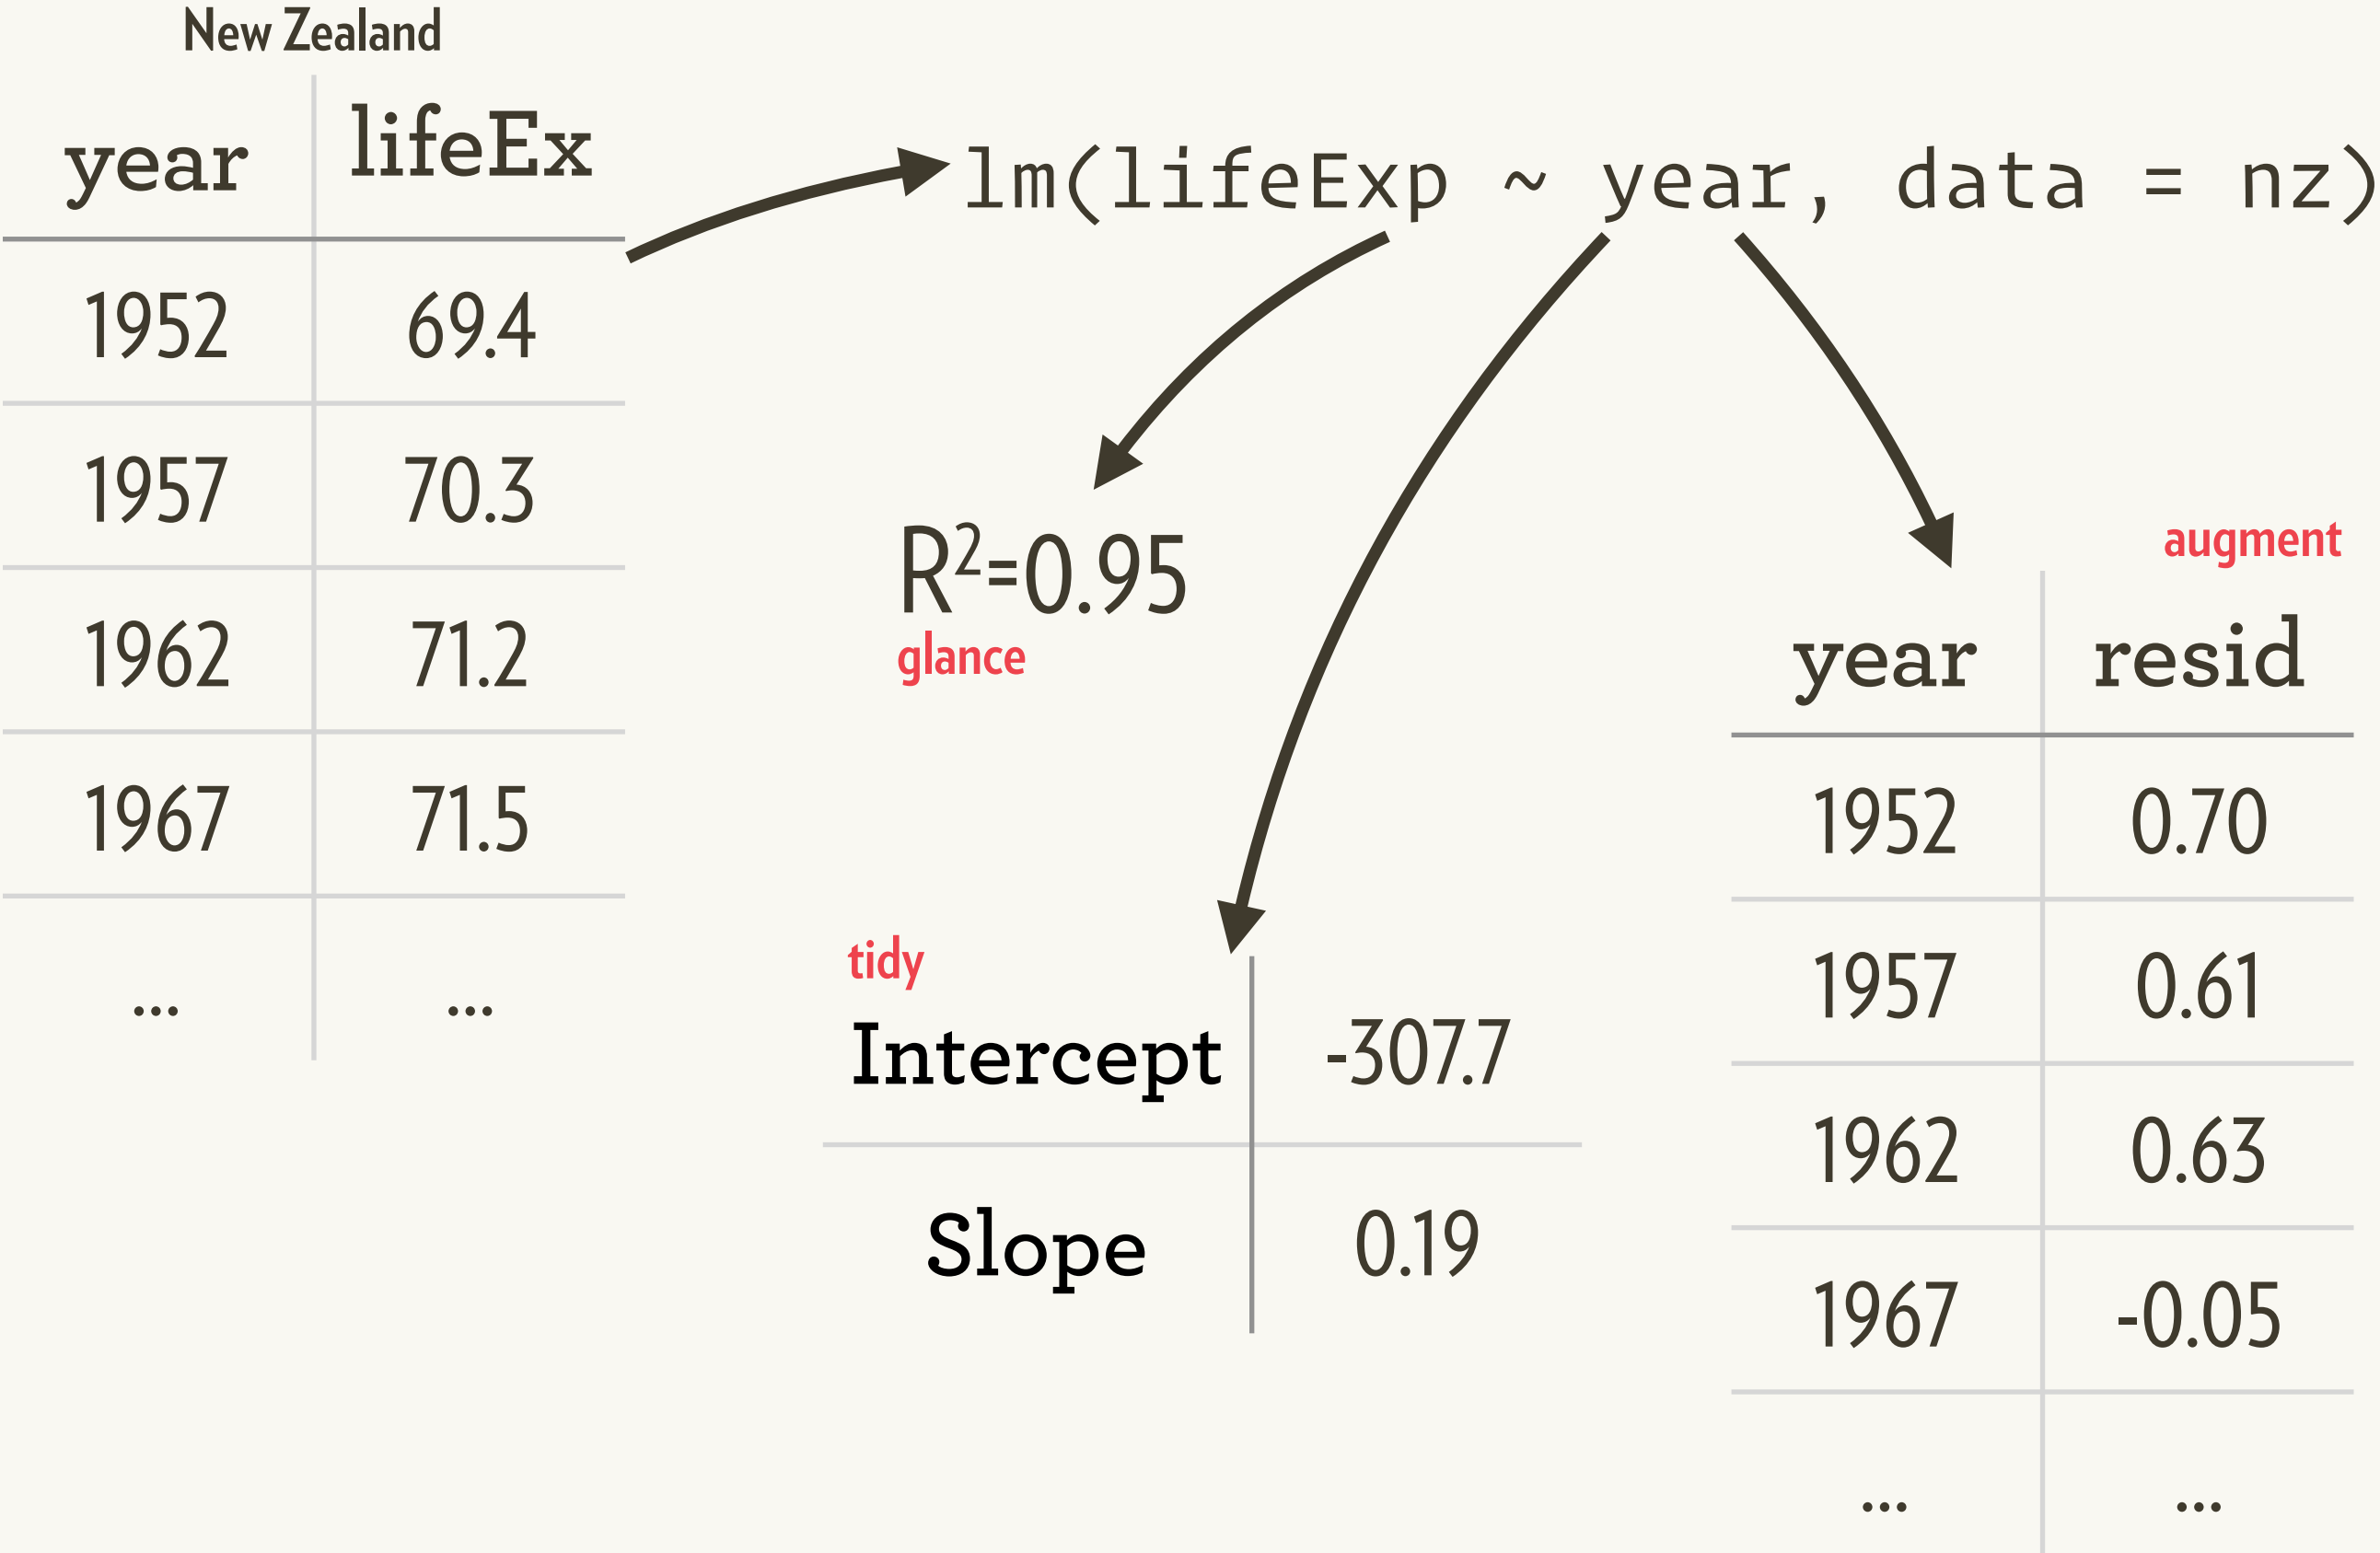
What can we do with a list of models?

Country	Data	Model
Afghanistan	<data>	<lm>
Albania	<data>	<lm>
Algeria	<data>	<lm>
...	<data>	<lm>

Models → tidy data

With broom, by David Robinson

What data can we extract from a model?



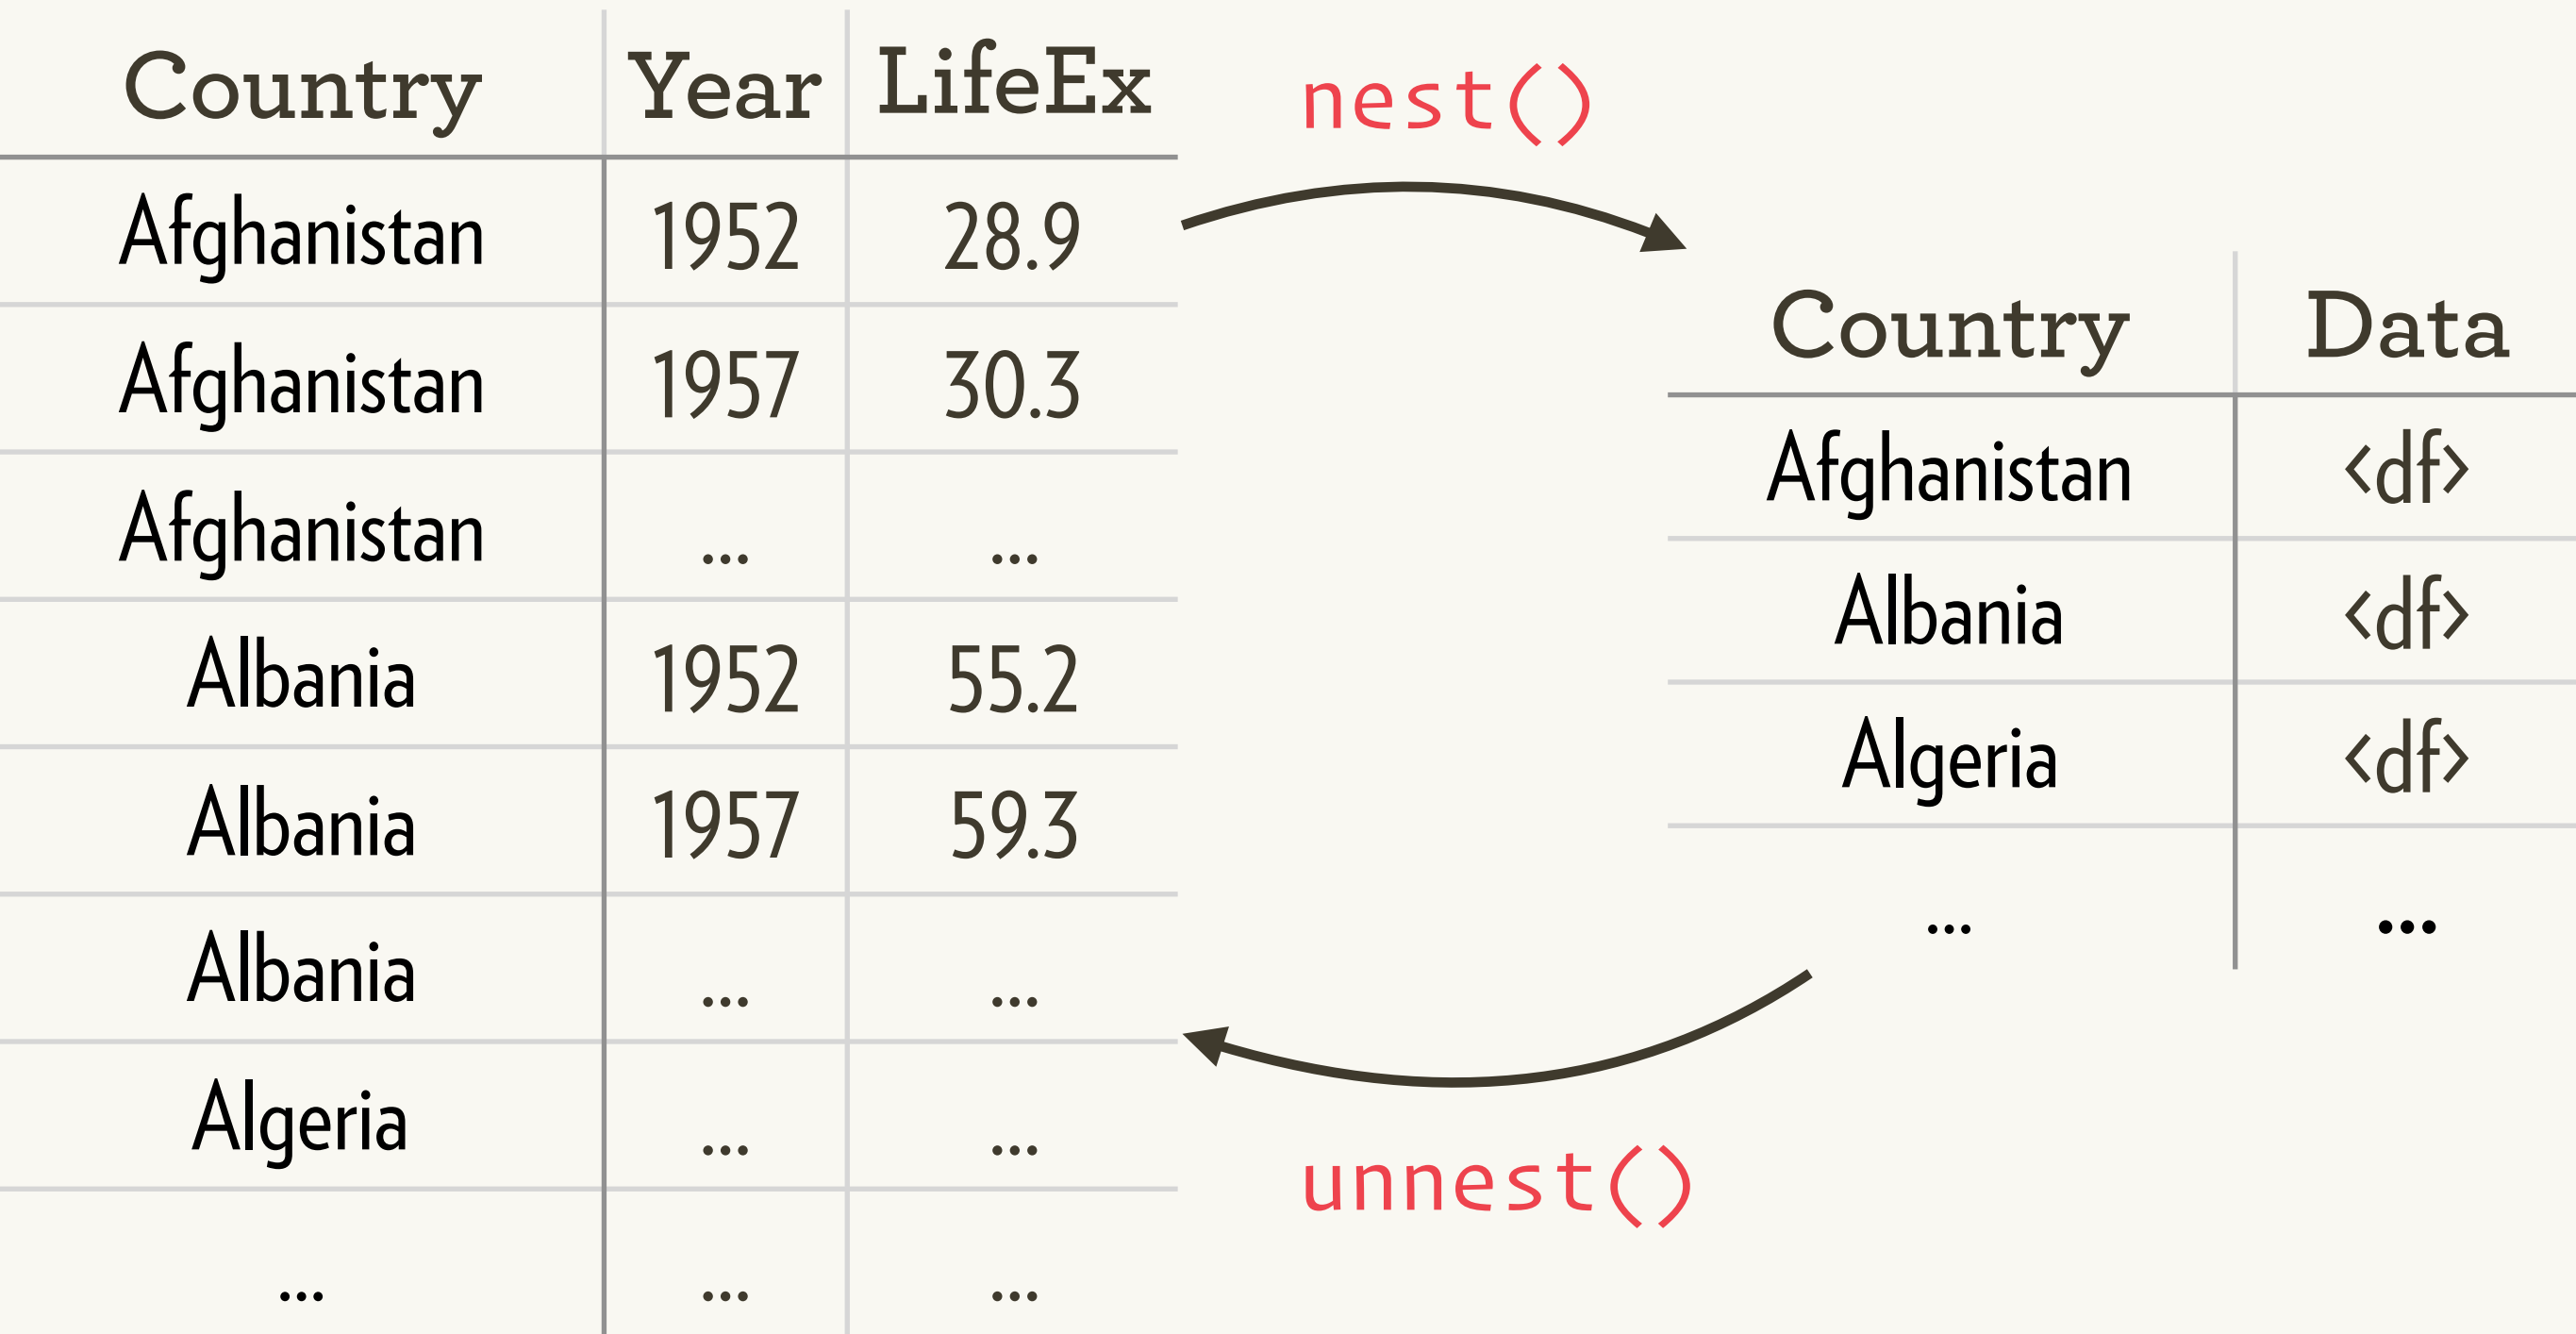
We need to do that for each model

```
models <- models %>%  
  mutate(  
    tidy      = map(model, broom::tidy),  
    glance    = map(model, broom::glance),  
    augment   = map(model, broom::augment)  
  )
```

Which gives us:

Country	Data	Model	Glance	Tidy	Augment
Afghanistan	<df>	<lm>	<df>	<df>	<df>
Albania	<df>	<lm>	<df>	<df>	<df>
Algeria	<df>	<lm>	<df>	<df>	<df>
...

Unnest lets us go back to a regular data frame




Demo

Conclusion

1. Store related objects in **list-columns**.
2. Learn **FP** so you can focus on verbs, not objects.
3. Use **broom** to convert models to tidy data.

dplyr
Data frames ← **broom** — Models



Lists
purrr

Workflow replaces many
uses of `ldply()/dplyr()` (**plyr**)
and `do() + rowwise()` (**dplyr**)

<http://r4ds.had.co.nz/>

This work is licensed under the
Creative Commons Attribution-Noncommercial 3.0
United States License.

To view a copy of this license, visit
<http://creativecommons.org/licenses/by-nc/3.0/us/>